

Dhruv Baslas

Data Engineer

📍 Noida, India | ✉️ dhruvbaslas@gmail.com | 📞 +91-8979217456 | 🔗 LinkedIn | 🐙 GitHub

PROFESSIONAL SUMMARY

Data Engineer with 6 years of experience designing scalable, metadata-driven ETL / ELT pipelines and Lakehouse architectures utilizing Azure, Databricks, and PySpark. Proven track record in implementing Medallion Architectures (Bronze / Silver / Gold), executing complex enterprise data migrations, and leading cross-functional engineering teams. Expert in building low-latency data ingestion frameworks that process high-volume datasets for global leaders in the financial, manufacturing, and healthcare sectors.

TECHNICAL SKILLS

Languages & Querying	Python, PySpark, SQL, Spark SQL, PL/SQL
Cloud & Platforms	Azure, Databricks, Unity Catalog, Azure Data Factory (ADF), ADLS, Key Vault, Logic Apps, Azure DevOps
Architecture	Medallion Architecture (Bronze/Silver/Gold), Lakehouse, Delta Lake, Metadata-Driven Pipelines, ETL/ELT
Orchestration & Sources	Databricks Workflows, Apache Airflow, AutoSys, SQL Server, Oracle, SAP S/4HANA, SAP BW

PROFESSIONAL EXPERIENCE

Senior Data Engineer <i>JD Software Pvt. Ltd.</i>	February 2025 – Present <i>Remote</i>
---	---

- Architected and deployed a Medallion (Bronze/Silver/Gold) Lakehouse in Azure Databricks, migrating legacy data warehouses to a metadata-driven framework and improving pipeline processing speed by 40%.
- Built event-driven ingestion pipelines using ADF and Azure Logic Apps to load multi-source enterprise data (SAP, Oracle, REST APIs) into ADLS, orchestrating scalable batch processing and Delta table creation via Databricks Workflows.
- Leveraged Unity Catalog across multiple Databricks workspaces to enforce centralized data governance, secure access controls, and manage data lineage for cross-functional enterprise analytics.
- Engineered complex PySpark and Spark SQL transformations to populate optimized Delta tables, supporting Slowly Changing Dimensions (SCD Type 2) historical tracking and reducing downstream dashboard query latency by 50%.
- Developed automated SQL and PySpark reconciliation frameworks across DEV, QA, and PROD environments, successfully catching 95% of data anomalies prior to production deployment.
- Led a team of engineers across multiple data domains and partnered with stakeholders to translate business requirements into production-ready pipelines, while mentoring junior staff on distributed system design.

Data Engineer – UnitedHealth Group <i>Accenture</i> <i>(Working as Team Lead)</i>	August 2024 – January 2025 <i>Bengaluru, India</i>
--	--

- Programmatically orchestrated complex data pipelines by integrating Apache Airflow with Databricks APIs to automate job submission and seamless workflow execution.
- Engineered a modern Medallion Architecture in Azure Databricks to process high-volume flat file data from ADLS into Delta Lake, enabling robust real-time analytics for over 50+ downstream consumers and decreasing data availability SLA times by 30%.

Data Engineer – UBS <i>Accenture</i> <i>(Progression: Analyst to Senior Software Engineer to Team Lead)</i>	September 2020 – July 2024 <i>Bengaluru, India</i>
--	--

- Architected end-to-end metadata-driven pipelines using Azure Data Factory (ADF), AutoSys, and PySpark to ingest and process over 500 GB of daily legacy mainframe and Oracle data into ADLS.
- Utilized Azure DevOps as the centralized Git repository for Databricks from 2020 onwards, rigorously managing branch policies, pull requests (PRs), and automated CI/CD pipelines to reduce deployment overhead by 50%.
- Developed custom PySpark exception-handling and validation logic to cleanse complex legacy data formats, enforcing strict data consistency and automated error reporting.

EDUCATION

Bachelor of Technology in Computer Science

GLA University

August 2016 – July 2020

Mathura, India

- **CGPA:** 8.00
- **Specialization:** Cloud Computing and Virtualization (in collaboration with IBM)

ACHIEVEMENTS & PUBLICATIONS

First Place Winner (IBM Hackathon)

Innovacer, IBM Bengaluru

- Ranked #1 for innovating and modernizing a legacy ATM system into a next-generation future system architecture.

Framework for Encryption and Authentication Cryptosystem

Publication Link

- Issue: Vol. 29 No. 10s (2020) - 4572-4576

CERTIFICATIONS

- **Microsoft Certified: Azure Data Engineer Associate (DP-203)**
- **Microsoft Certified: Azure Data Fundamentals (DP-900)**
- **Microsoft Certified: Azure Fundamentals (AZ-900)**
- **Microsoft Certified: Power Platform Fundamentals (PL-900)**
- **IBM Certified: Cloud Computing & Virtualization Graduate**